

Gustavo Balbi Queiroz

Analytics Engineer · Data Modeling · dbt · SQL · BI

gustavobalbi03@gmail.com | +55 (91) 98208-4478 | linkedin.com/in/gustavo-balbi | github.com/gustavobalbi

SUMMARY

Analytics engineer focused on the transformation layer — turning raw data into clean, tested, business-ready models. Dimensional modeling (star schema) and analytical SQL (CTEs, window functions, JSON) on a daily basis, plus delivery in Power BI for executive consumption, at a large company in the automotive sector. Requirements gathering with business areas and translation into governed metrics (RLS). Personal project on a modern stack (dbt + Apache Airflow on Azure) with automated data-quality tests. Background in Python and applied research in neural networks. Fluent in English.

EXPERIENCE

Data Analyst — Revemar Gestão de Negócios

Mar/2026 – Present

Full data lifecycle across an end-to-end pipeline: extraction from Oracle, transformation in Apache Hop, loading into SQL Server under a Medallion architecture, and delivery for consumption in Power BI, with loads orchestrated via Jenkins.

- **Data modeling & analytical SQL** — I do dimensional modeling (facts and dimensions) and write analytical SQL daily: CTEs, window functions, and handling of semi-structured data (JSON).
- **Power BI** — built dashboards for executive consumption — parts inventory, vehicle inventory, report access and permissions, sales performance, and financial performance.
- **Requirements gathering** — led requirements gathering with business areas across different segments to design new data projects; conducted technical interviews when selecting new team members.
- **Data access governance** — implement RLS (row-level security) — an environment with 52 users across multiple permission levels.
- **End-to-end ETL/ELT pipelines** — I build pipelines under a Medallion architecture, extracting from Oracle and loading into SQL Server, through to delivery for analytical consumption.
- **Performance optimization (recurring)** — I continuously optimize queries and ETLs: cut jobs from ~10 to ~5 min (ending queues and executor overload); in another case, reduced a load from hours to seconds by eliminating a cartesian product and creating indexes.
- **AI adoption in the department** — led the integration of AI assistants into the data environment via MCP (Power BI, Apache Hop, database) and built the business case (costs, benefits, and risks) for the corporate subscription.
- **Orchestration observability** — developed a monitoring solution pulling Jenkins metrics via API (duration, queue, triggers, executor utilization, errors) into a dashboard that gives the team performance visibility and non-technical management transparency into the environment.
- **Database maintenance & administration** — maintain a portfolio of 70+ reports in production (enhancements, fixes, and recovery of ETLs disrupted by business-rule changes) and administer the database server — storage, performance, schemas, and versioning.
- **Automation (RPA)** — automated report distribution via WhatsApp delivery — a first-of-its-kind solution at the company.

Undergraduate Researcher — Analytical Intelligence Lab, UFPA

2025 – Present

- Supervised classification models (scikit-learn) to predict wildfire risk in the Legal Amazon, comparing an MLP against a Random Forest baseline.
- Owner of the project pipeline — preprocessing, experimentation, and evaluation; the MLP reached ~95% accuracy (macro-F1 0.94), outperforming the baseline.

Django Developer — ASPARÁ (Public Servants Association)

Dec/2025 – Mar/2026

- Institutional web system in Django (Python) deployed to production on Azure cloud, with an authenticated area and different access levels.
- Relational database modeling and integration in SQL Server for form records and user/permission control.
- Exploratory analysis of submitted forms (Pandas, Plotly) to process and visualize the information.

Teaching Assistant — Introductory Programming, UFPA (2025): taught algorithms and programming logic in Python to incoming Computing freshmen.

PROJECTS

High-Frequency Data Pipeline (personal project)

- **End-to-end ELT** — ingestion from a high-frequency public API (real-time public-transit positions) → Azure SQL → layered transformation with dbt (Medallion) → orchestration via Apache Airflow (Docker) → consumption layer in Power BI.
- **SQL modeling** — handling semi-structured JSON with schema-on-read in the bronze layer (OPENJSON) and CTE-based transformations through to the consumption models.
- **Data quality** — automated tests with dbt (not_null, unique, relationships) ensuring reliability across layers.

Fuzzy System for Technical-Ticket Prioritization (academic project)

- **Fuzzy inference (Mamdani)** — recommends ticket priority from service impact, percentage of affected users, and time open, with an output from 0 to 100 (low/medium/high/critical).
- **Engineering and testing** — manual implementation in Python validated by cross-checking against scikit-fuzzy, packaged into modules with around 30 automated tests.

SKILLS

Languages: Python, SQL, DAX

Modeling & Data Warehouse: dimensional modeling (star schema — facts and dimensions), Medallion architecture, RLS (row-level security), semi-structured data / JSON

Data Transformation & Quality: dbt (not_null / unique / relationships tests), SQL (CTEs, window functions), Power Query

BI & Visualization: Power BI, DAX, Pandas, Plotly

Orchestration & Pipelines: Apache Airflow, Jenkins, Apache Hop (Pentaho/PDI-family ETL), dbt

Databases: SQL Server, Oracle

Cloud & DevOps: Azure, Docker, Git

EDUCATION

Information Systems — Federal University of Pará (UFPA). Expected graduation: 2027.

QUALIFICATIONS & CERTIFICATIONS

- Artificial Neural Networks in Python — Udemy (2025)
- Python Fundamentals — DIO (2024) · Getting Started with SQL — DIO (2025)
- Web Development with Python and Django — DIO (2025) · HTML5 and CSS3 — DIO (2022)
- English (fluent) — Brazil–United States Binational Center (2013–2020)

LANGUAGES

Portuguese (native) · English (fluent) · Spanish (intermediate)